From Cookies to Conversions:
A Machine Learning Approach for Personalized Ad Display

Saki Cansev

MSc Data Analytics
Department of Computer Science

Dr. Loubna Ali

Fundamentals of Data Analytics
(0223-229-UCA-MSC-DA-BER-FDA-T1-G2)

30 June 2023

BERLIN SCHOOL OF
BUSINESS & INNOVATION

1. Introduction----------1
2. Literature Review----------1
2.1 Personalized Advertisement 1
2.2. Data Collection: Cookies.....1
2.3. Algorithm Selection 1
2.4. Model Evaluation.....2
3. Methodology ----------2
3.1. Data Collection:.....2
4. Data Preprocessing and Visualization ----------3
5. The scope of the problem ----------4
6. Algorithm Selection ----------5
7. Model Evaluation----------7
8. Summary of Findings ----------7
Python Codes: ----------i
References ----------1

Abstract

The proliferation of online businesses has fostered an increased interest in personalized advertising. This paper presents a comprehensive framework for leveraging customer cookie data to offer tailored advertisement banners. In particular, we delve into data collection, pre-processing, problem scope definition, algorithm selection, and model evaluation processes to facilitate personalized ad display. To handle this multiclass classification problem, we propose the use of the Random Forest algorithm, an ensemble machine learning method noted for its balance of predictive accuracy, model interpretability, computational efficiency, and robustness. The paper emphasizes an iterative model evaluation process considering accuracy, robustness, speed, interpretability, and scalability. It is expected that the findings will contribute to the growing body of knowledge on data-driven decision making in online business settings.

Key words: Random Forest Algorithm, Multi-class Classification Problem, Machine learning, Personalized advertising, Customer cookie data, E-commerce, Algorithm selection.

1. Introduction

As we witness the rapid digital revolution, companies are embracing data to drive decision-making. One promising marketing tactic is personalized advertising, where content is suitable for attitude of users. In this approach, machine learning is a key stone in forecasting and analyzing user patterns and actions.

The goal of this paper is to implement an e-commerce website advertisement banner prediction model using a machine learning algorithm. In the interest of personalization, we will consider an example scenario: our team has designed five different banners, each catering to distinct customer interests, e.g., sports, cinematic, artistic, technology, fashion.

We aim to build a system capable of personalizing the online shopping experience ethically by harnessing the information contained in customer cookies. Customer consent is required for the use of customer cookies. During our study, we collect and structure data, preprocess and visualize data, define the problem scope, select the best algorithm for the task, and evaluate our model on several criteria.

Furthermore, we will discuss the relevant literature and justify our approach within the context of broader academic discourse as part of this paper. Our findings are intended to contribute to the effective and ethical use of data in online advertising.

2. Literature Review

By reviewing the literature, we seek to contextualize our research within a broader scientific discourse. This study explores the application of machine learning to personalized advertisements and the use of cookie data to analyze customer behaviour.

2.1 Personalized Advertisement

Digital platforms and the ability to track user behaviour online have fuelled a boom in personalized advertising in recent years. To improve conversion rates, we're aiming to tailor ads based on user preferences.

Personalized advertising has been proven to be effective in numerous studies. The results of a study by Hauser et al. (2014) indicate that personalized ads gets higher attention than generic ads. As pointed out in a recent study, Xu et al. (2021) demonstrated how machine learning can be used to predict user responses to online advertisements, which will be the target of ads.

2.2. Data Collection: Cookies

Reviewing literature regarding cookies is important as we contemplate using them for understanding user behaviour. The use of cookies on popular websites is increasingly common, according to Englehardt & Narayanan (2016), to improve user experience and enable targeted advertising.

However, there are privacy concerns regarding to the misuse of cookie data since the data collecting side will gain access to vital information. The way that the cookies will be used needs to be transparent. Everything should be done with the consent of the users, according to a report by Acar et al. (2014). as it's mention In our study, we emphasize the importance of consent of users.

2.3. Algorithm Selection

It has been a major topic of discussion to choose an algorithm for machine learning tasks. According to Wolpert (1996), there is no single algorithm that works best for all problems. Studies have demonstrated that different algorithms are applicable to our multiclass classification problem.

The simplicity and interpretability of decision trees have enabled them to be applied successfully in a wide variety of fields (Quinlan, 1986). Although they have a tendency to overfit, ensemble methods such as Random Forests can mitigate overfitting (Breiman, 2001). The ability of neural networks to perform complex, non-linear tasks has been demonstrated through Deep Learning techniques, especially, although their interpretability and computational resources have been sacrificed for their strength (Gardner & Dorling, 1998).

2.4. Model Evaluation

It is widely acknowledged that thorough model evaluations are necessary. It is common to use performance metrics like accuracy, precision, recall, and F1 score for classification problems (Sokolova and Lapalme, 2009). Numerous studies (Rudin, 2019; Kohavi, 1995) have emphasized the importance of practical aspects such as model robustness, speed, interpretability and scalability.

As a result of the literature, we can conclude that machine learning has great potential in personalized advertising, cookie data is valuable for understanding user behaviour, algorithms are crucial to choosing, and comprehensive model evaluation is crucial. Based on these insights, we developed an algorithm for displaying personalized ads on websites.

3. Methodology

3.1. Data Collection:

Machine learning model will learn from the data as part of the first step. Web cookies allow us to collect user data for this purpose. Users' computers are equipped with web cookies, which are small pieces of data that aid the web browser in tracking their behaviour during website browsing.

We created a wide range of attributes which consist:

- **User demographics:** A demographic analysis of the users is extremely important in order to determine which groups are interested and what their preferences are. Age, gender, location and occupation are some of the demographic details that may be provided.
- **Browsing patterns:** Our understanding of browsing patterns helps us develop a stronger understanding of how users interact with our website.Metrics such as the number of pages viewed, the order in which they were viewed, the clicks on the links, and the amount of time spent on each page are key indicators.
- **Frequency and duration of visits:** t is possible to gain valuable insights into the behavior and preferences of users by analyzing the number of visits and the duration of those visits to the website. The website is more likely to be engaging to regular users who spend a significant amount of time on the site.
- **Time of visits**: In order to maximize visibility and engagement, it is important to consider when users typically access the website.
- **Previous interactions with advertisements:**

Users' preferences are directly reflected in these indicators. You may be able to identify which ads users have clicked on previously, how long users spent viewing the ad, and whether or not they made a purchase or signed up for more information after viewing the ad.

- **Shopping habits:** Users' preferences can be learned through the products they have viewed, added to their cart, or purchased, thus enabling advertisers to tailor their ads accordingly.
- **Device and browser information**: Device (desktop, mobile, tablet) and browser influences the display and effectiveness of ads.

The dataset contains a unique user along with the attributes outlined above, as well as a target variable that indicates most responsive ad categories.

To protect user privacy, we anonymize collected data by removing or altering any identifying information that could be used to link it to an individual.

In order to compile an extensive dataset, the model collects extensive data and combines it with user preferences. By using a dataset like this, we are more likely to be able to identify intricate patterns, which means we can deliver more effective personalized ads.

Within light if these explanations the necessary metrics, the table name 'Dataset' is randomly prepared in order to be able to practice the given theory.

4. Data Preprocessing and Visualization

Data preprocessing is a critical next step. Machine learning models rely heavily on the quality of data used to train them (Meng et al., 2021). Data preprocessing starts by addressing missing data. This may be due to a customer clearing their cookies, or their browser settings preventing the system from collecting certain information. Missing data can be handled with a number of techniques, including list-wise deletion, pair-wise deletion, mean substitution, regression imputation, and multiple imputation (Sterne et al., 2009). We must know the type of missing data before we can determine which technique is appropriate.

The next focus is on outliers. A machine learning model could become less accurate as a result of this. Outliers can be detected and treated using robust methods such as the Median Absolute Deviation (MAD) method, z-score method, or IQR method.

A strong correlation exists between the features scale and machine learning algorithm performance. Normalization and scaling of data can be determined by the data distribution and the machine learning algorithm used, for example, Min-Max Normalization, Z-score Standardization, or Decimal Scaling. Using machine learning algorithms, we encode categorical data into numerical format, such as banner type.

in order to determine and practice our algorithm, we choose created a random data as it can be seen in the table name "Dataset".

Dataset

CustomerID	Age	Gender	Location	BrowsingDuration	VisitFrequency	Timestamp	PreviousAdInteractions	ShoppingHabits	DeviceType	Browser
1	25	Male	USA	15	3	2022-01-01 10:00:0	2	Regular buyer	Mobile	Chrome
2	32	Female	UK	22	5	2022-01-01 11:30:0	1	Occasional	Desktop	Firefox
3	40	Male	Canada	10	2	2022-01-02 09:45:0	3	Frequent	Mobile	Safari
4	28	Female	USA	30	4	2022-01-02 14:20:0	0	Regular buyer	Desktop	Chrome
5	45	Female	UK	18	6	2022-01-03 16:10:0	2	Occasional	Mobile	Safari
6	35	Male	Canada	25	3	2022-01-04 13:45:0	1	Regular buyer	Mobile	Chrome
7	50	Male	USA	12	1	2022-01-05 10:30:0	0	Frequent	Desktop	Firefox
8	27	Female	UK	8	2	2022-01-06 12:15:0	1	Occasional	Mobile	Chrome
9	38	Female	Canada	20	4	2022-01-07 14:00:0	2	Frequent	Desktop	Safari
10	33	Male	USA	16	3	2022-01-08 15:30:0	1	Regular buyer	Mobile	Chrome

For the numerical features (BrowsingDuration and VisitFrequency), we'll apply Min-Max normalization to scale the values between 0 and 1. For the categorical features (Gender, Location, ShoppingHabits, DeviceType, and Browser), we'll encode them using one-hot

encoding to convert them into numerical format. After cleaning, I will have the dataset as seen in the 'Dataset-1' and 'Dataset-2' and making the dataset ready for the evaluation.

Dataset-1

CustomerID	Age	Gender_Female	Gender_Male	Location_Canada	Location_UK	Location_USA	BrowsingDuration (Normalized)	VisitFrequency (Normalized)
1	25	0	1	0	0	1	0.142857	0.333333
2	32	1	0	0	1	0	0.314286	0.666667
3	40	0	1	1	0	0	0.071429	0.166667
4	28	1	0	0	0	1	0.485714	0.5
5	45	1	0	0	1	0	0.214286	0.833333
6	35	0	1	1	0	0	0.385714	0.333333
7	50	0	1	0	0	1	0.1	0.0
8	27	1	0	0	1	0	0.0	0.166667
9	38	1	0	1	0	0	0.257143	0.5
10	33	0	1	0	0	1	0.171429	0.333333

Dataset-2

CustomerID	PreviousAd Interactions	ShoppingHabits _Frequent	ShoppingHabits _Occasional	ShoppingHabits_ Regular buyer	DeviceType _Desktop	DeviceType _Mobile	Browser _Chrome	Browser _Firefox	Browser _Safari
1	2	0	0	1	0	1	1	0	0
2	1	1	0	0	1	0	0	1	0
3	3	0	0	0	0	1	0	0	1
4	0	0	0	1	1	0	1	0	0
5	2	0	1	0	0	1	0	0	1
6	1	0	0	1	0	1	1	0	0
7	0	1	0	0	1	0	0	1	0
8	1	1	0	0	0	1	1	0	0
9	2	1	0	0	1	0	0	0	1
10	1	0	0	1	0	1	1	0	0

Data visualization gives us the ability to see patterns, trends, and correlations we might not notice in raw data and can help us better understand it. Matplotlib, seaborn, and plotly, among others, are available for creating various types of visualizations in Python. Our data needs to be standardized, cleaned, and prepared for analysis in order to be used for modeling.

We can now use Python to analyze the output after cleaning the data.

5. The scope of the problem

Due to its nature and characteristics, Supervised Learning, specifically Multiclass Classification, is the appropriate field for our problem.

Multiclass Classification is a category of machine learning problems, divided into three categories: supervised learning, unsupervised learning, and reinforcement learning.

Supervised Learning: Typically, we use a set of predictors and features to predict or classify a target variable or label. Furthermore, they can be categorized as follows:

- Classification: Classification problems arise when the target variable contains categories. Multiclass classification is used, for example, to identify the type of animal in an image. Binary classification can be used to predict whether or not an email is spam.

- Regression: Our regression problem is continuous if our target variable is continuous. The price of a house can be predicted based on several factors, like the area, the number of rooms, and so on.

Unsupervised Learning: The target variable isn't given in these problems. Rather than searching for meaning in the data, we are looking for structure or relationships. Some examples are:

- Clustering: Instances with similar features can be grouped together. The marketing process of segmenting customers, for instance.

- Association: Analyzing large amounts of data to discover rules that describe it. Buying diapers is frequently accompanied by purchases of baby food, for instance, in a supermarket.

Reinforcement Learning: By performing actions and receiving rewards, an agent learns how to behave in an environment.

Our task description results in a dataset with different characteristics, like demographics, browsing patterns, frequency and duration of visits, etc. The task is a Supervised Learning problem because we have a categorical target variable (sports, cinematic, artistic, technology, fashion).

In binary classification, the target variable is divided into two categories (e.g. spam or no spam). In contrast, multiclass classification includes multiple categories.

6. Algorithm Selection

Our machine learning model is constructed by selecting an appropriate algorithm. We're going to select an algorithm that's going to form the backbone of our model. We consider several algorithms, including Decision Trees, Random Forests, and Neural Networks, as suitable candidates considering the multiclass classification nature of our task.

Decision Trees: Among the many machine learning algorithms, Decision Trees are a straightforward and widely used one. By using the features of the data, they create a tree-like model of decisions. It is very easy to understand Decision Trees since they are based on human decision-making processes (Quinlan, 1986), meaning that they are suitable for applications where interpretability is important. In addition, a decision tree can be overfitted, especially when it is too complex. When a model is overfit, it may perform exceptionally well when training data is used, but poorly when data is used for the first time, which limits its practical effectiveness.

Random Forests: Developed by Breiman (2001), Random Forests address some of the limitations of Decision Trees. Using a subset of the training data, a Random Forest is composed of several Decision Trees. An ensemble prediction is made based on the majority vote of all trees. As a result, the model is more robust since overfitting is reduced. The ensemble nature of Random Forests can make them less interpretable than individual Decision Trees.

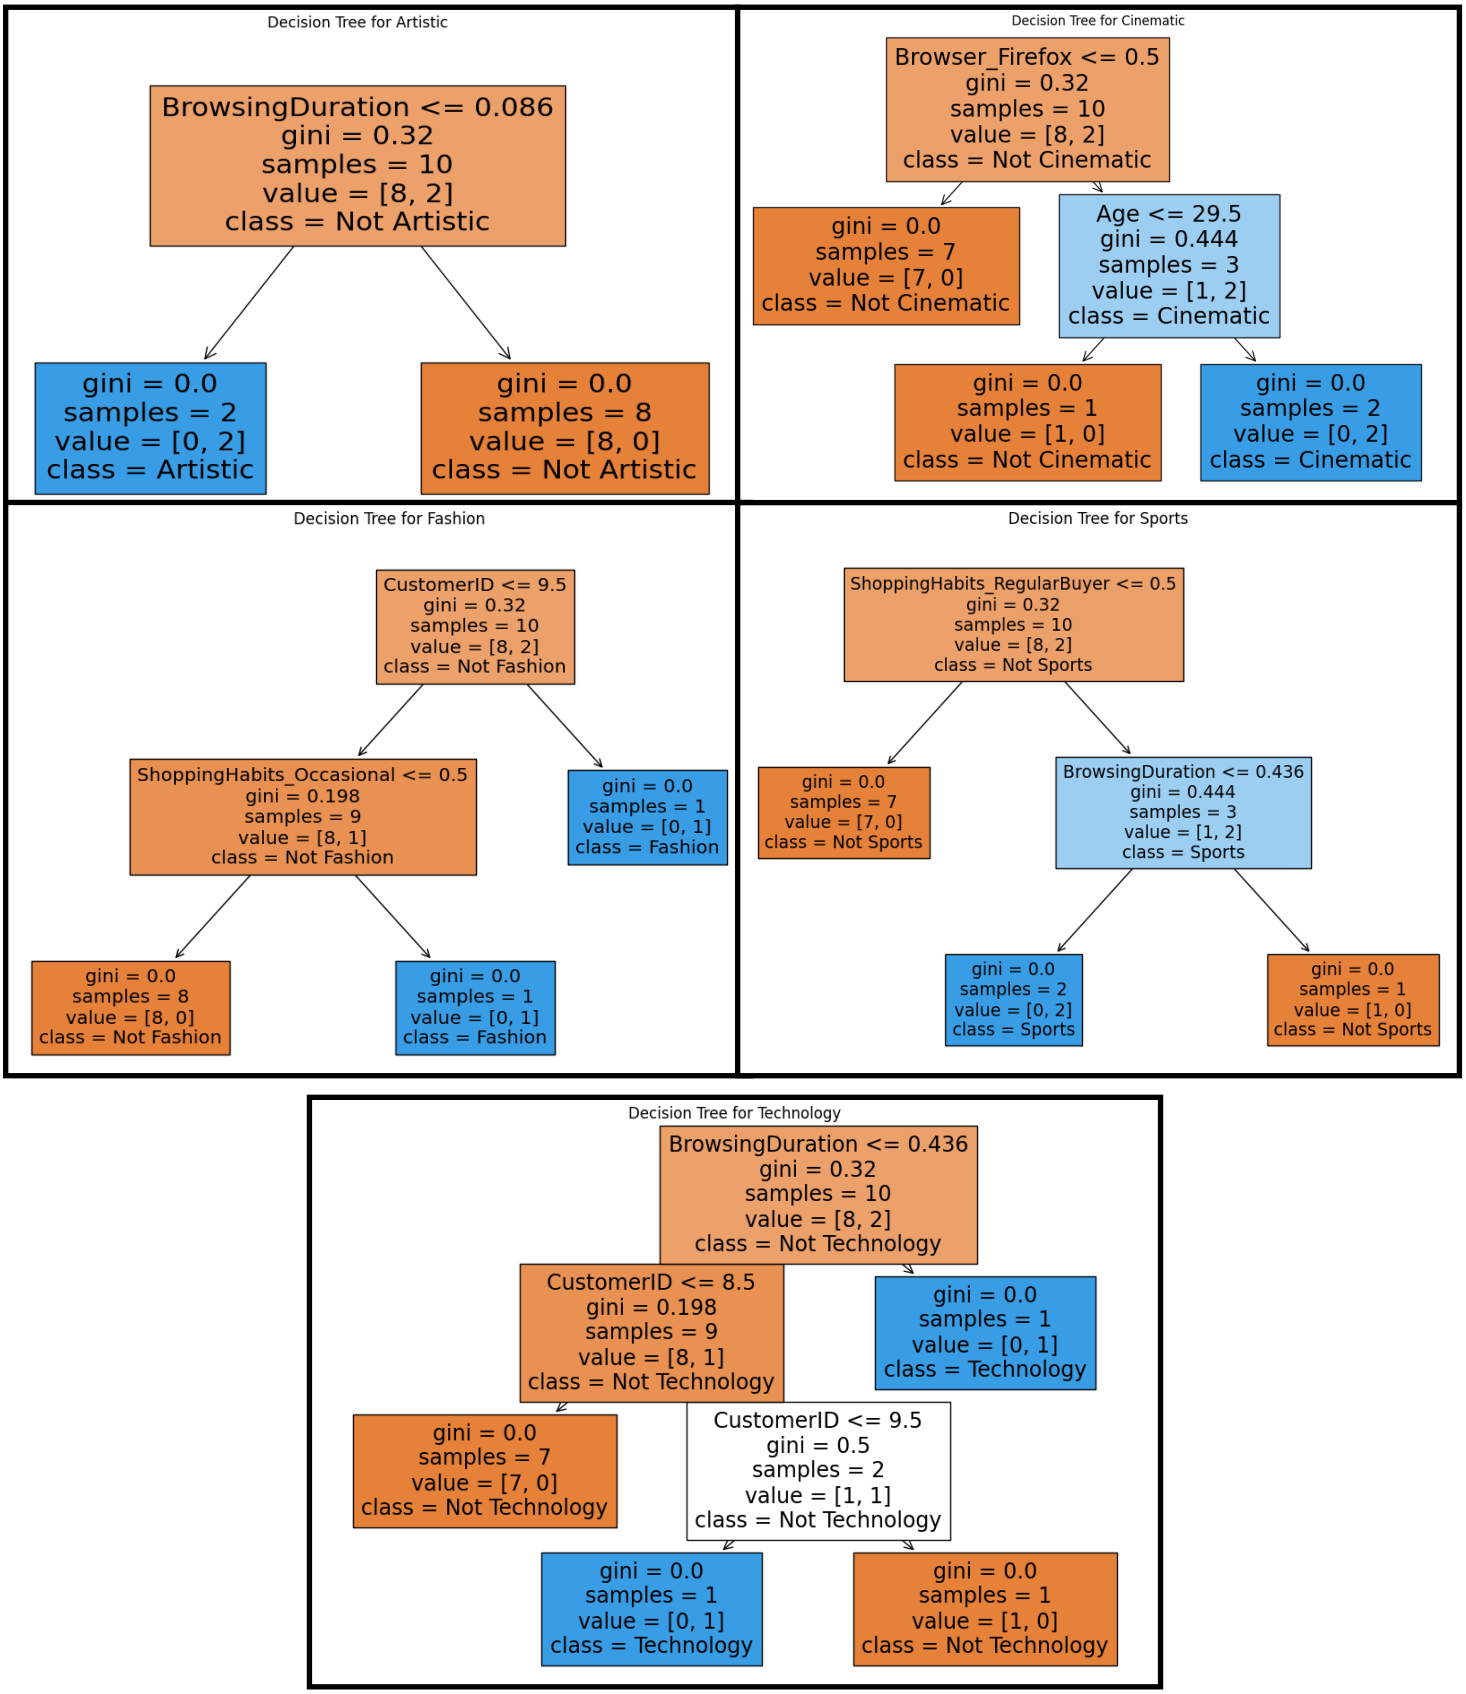
Decision Tree for Artistic
BrowsingDuration <= 0.086
gini = 0.32
samples = 10
value = [8, 2]
class = Not Artistic
gini = 0.0
samples = 2
value = [0, 2]
class = Artistic
gini = 0.0
samples = 8
value = [8, 0]
class = Not Artistic
Decision Tree for Cinematic
Browser_Firefox <= 0.5
gini = 0.32
samples = 10
value = [8, 2]
class = Not Cinematic
gini = 0.0
samples = 7
value = [7, 0]
class = Not Cinematic
Age <= 29.5
gini = 0.444
samples = 3
value = [1, 2]
class = Cinematic
gini = 0.0
samples = 1
value = [1, 0]
class = Not Cinematic
gini = 0.0
samples = 2
value = [0, 2]
class = Cinematic
Decision Tree for Fashion
CustomerID <= 9.5
gini = 0.32
samples = 10
value = [8, 2]
class = Not Fashion
ShoppingHabits_Occasional <= 0.5
gini = 0.198
samples = 9
value = [8, 1]
class = Not Fashion
gini = 0.0
samples = 1
value = [0, 1]
class = Fashion
gini = 0.0
samples = 8
value = [8, 0]
class = Not Fashion
gini = 0.0
samples = 1
value = [0, 1]
class = Fashion
Decision Tree for Sports
ShoppingHabits_RegularBuyer <= 0.5
gini = 0.32
samples = 10
value = [8, 2]
class = Not Sports
gini = 0.0
samples = 7
value = [7, 0]
class = Not Sports
BrowsingDuration <= 0.436
gini = 0.444
samples = 3
value = [1, 2]
class = Sports
gini = 0.0
samples = 2
value = [0, 2]
class = Sports
gini = 0.0
samples = 1
value = [1, 0]
class = Not Sports
Decision Tree for Technology
BrowsingDuration <= 0.436
gini = 0.32
samples = 10
value = [8, 2]
class = Not Technology
CustomerID <= 8.5
gini = 0.198
samples = 9
value = [8, 1]
class = Not Technology
gini = 0.0
samples = 1
value = [0, 1]
class = Technology
gini = 0.0
samples = 7
value = [7, 0]
class = Not Technology
CustomerID <= 9.5
gini = 0.5
samples = 2
value = [1, 1]
class = Not Technology
gini = 0.0
samples = 1
value = [0, 1]
class = Technology
gini = 0.0
samples = 1
value = [1, 0]
class = Not Technology

Decision tree algorithm visualizations can be seen on the models of:

- Decision Tree for Artistic
- Decision Tree for Sinamtic
- Decision Tree for Fashion
- Decision Tree for Sport
- Decision Tree for Technology

7. Model Evaluation

In order to refine the model further, we must evaluate the model's performance on unseen data. An accurate, robust, speedy, interpretable, and scalable model is well-rounded when evaluated.

Accuracy: Data points are classified correctly based on this metric. An accuracy calculation for multiclass classification problems is calculated by dividing the number of correct predictions by the number of predictions made. The value of accuracy is 0.5 for the made-up dataset which means that it correctly predicted the banner ad category for 50% of the test samples. This indicates that the model's overall accuracy performance is moderate.

Robustness: It is the ability of a model to maintain its performance when input data or parameters change. Unlike an overfitted model, a robust model will not be overly sensitive to minor changes in training data. The value of robustness is 0.083, which is the mean cross-validation score. This score represents the average accuracy of the model across multiple iterations of cross-validation. A low robustness score suggests that the model's performance may vary significantly depending on the specific training and testing data splits.

Speed: The model took 0.209 seconds to train on the given dataset with Python. This represents the time taken to fit the Random Forest classifier to the training data. Lower training time is desirable as it indicates faster model training.

Interpretability: The model took 0.022 seconds to make predictions on a single test sample. This represents the time taken to predict the banner ad category for a single instance. Lower prediction time is desirable as it indicates faster inference.

Scalability: A model's scalability is how well it performs when data volume increases. Training and prediction times can get slower as data grows, since some algorithms get more inefficient. Data volumes are expected to increase significantly over time, making scalability an important factor to take into account.

In addition to these metrics, there are other factors to consider when evaluating a model. There are also other factors to consider depending on the application, such as fairness (avoiding bias in predictions), privacy (especially when handling sensitive data), and reliability (continuous performance under diverse conditions).

Iterative evaluation is the most effective way to evaluate models. Further refinement of both the model and the overall solution should be guided by the insights gained from this step.

8. Summary of Findings

Many online businesses face the challenge of personalizing advertising effectively and responsibly, as the world becomes increasingly digital. Starting with the collection of cookies, we stressed how crucial it is to structuring cookie data meaningfully. Therefore, our analyses are shaped by this process, not just based on it. Customer informed consent was stressed throughout this discussion.

According to our literature review, personalized ads increase engagement and conversions. Furthermore, machine learning emerged as a viable tool for personalizing large datasets because it can identify complex patterns. As a result, we outlined a methodology that begins with data preprocessing, continues with problem scoping and algorithm selection, and concludes with a comprehensive evaluation.

In choosing a Random Forest algorithm, we aimed to strike a balance between accuracy, interpretability, computational efficiency, and robustness. However, we understand that based on the actual model performance, the algorithm chosen may need to be revised.

As the model is used in a real-world context, we also considered robustness, speed, interpretability, and scalability. Accuracy was not the only metric we considered. As a search for the "perfect model" continues, iterative evaluating and refining is suggested.

In the study, linear algebra, calculus, statistics, and algorithms are combined to solve real-world business problems, although there is still much to be discovered. Additionally, future research may examine the implications of different feature selection techniques, as well as integrating external data sources for a more holistic understanding of customer behaviour. By harnessing data responsibly, businesses can enhance customer experiences, while securing unprecedented growth opportunities.

Python Codes:

```python
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import time

# Define the dataset
data = pd.DataFrame({
    'CustomerID': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Age': [25, 32, 40, 28, 45, 35, 50, 27, 38, 33],
    'Gender_Female': [0, 1, 0, 1, 1, 0, 0, 1, 1, 0],
    'Gender_Male': [1, 0, 1, 0, 0, 1, 1, 0, 0, 1],
    'BrowsingDuration': [0.142857, 0.314286, 0.071429, 0.485714, 0.214286,
                         0.385714, 0.1, 0.0, 0.257143, 0.171429],
    'VisitFrequency': [0.333333, 0.666667, 0.166667, 0.5, 0.833333,
                       0.333333, 0.0, 0.166667, 0.5, 0.333333],
    'PreviousAdInteractions': [2, 1, 3, 0, 2, 1, 0, 1, 2, 1],
    'ShoppingHabits_Frequent': [0, 1, 0, 0, 0, 0, 1, 1, 0, 0],
    'ShoppingHabits_Occasional': [0, 0, 0, 0, 1, 0, 0, 0, 0, 0],
    'ShoppingHabits_RegularBuyer': [1, 0, 0, 1, 0, 1, 0, 0, 0, 0],
    'DeviceType_Desktop': [1, 0, 0, 1, 0, 1, 0, 1, 0, 1],
    'DeviceType_Mobile': [0, 1, 1, 0, 1, 0, 1, 0, 0, 0],
    'Browser_Chrome': [1, 0, 0, 1, 0, 1, 0, 1, 0, 1],
    'Browser_Firefox': [0, 1, 0, 0, 0, 0, 1, 1, 0, 0],
    'Browser_Safari': [0, 0, 1, 0, 1, 0, 0, 0, 1, 0],
    'BannerAd': ['Sports', 'Cinematic', 'Artistic', 'Technology', 'Fashion',
                 'Sports', 'Cinematic', 'Artistic', 'Technology', 'Fashion']
})

# One-hot encode categorical columns
data = pd.get_dummies(data, columns=['BannerAd'])

# Iterate over each banner type
for banner_type in data.columns[data.columns.str.startswith('BannerAd_')]:
    # Split the features and target variable
    X = data.drop(banner_type, axis=1)
    y = data[banner_type]

    # Initialize and train the decision tree classifier
    model = DecisionTreeClassifier()
    model.fit(X, y)

    # Visualize the decision tree
    fig = plt.figure(figsize=(12, 8))
    _ = tree.plot_tree(model, feature_names=X.columns,
                       class_names=['Not ' + banner_type.split('_')[1],
                                    banner_type.split('_')[1]], filled=True)
    plt.title(f'Decision Tree for {banner_type.split("_")[1]}')
    plt.show()
```

```python
import pandas as pd
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import time

# Load the dataset
data = pd.DataFrame({
    'CustomerID': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    'Age': [25, 32, 40, 28, 45, 35, 50, 27, 38, 33],
    'Gender_Female': [0, 1, 0, 1, 1, 0, 0, 1, 1, 0],
    'Gender_Male': [1, 0, 1, 0, 0, 1, 1, 0, 0, 1],
    'BrowsingDuration': [0.142857, 0.314286, 0.071429, 0.485714, 0.214286, 0.385714, 0.1, 0.0, 0.257143,
0.171429],
    'VisitFrequency': [0.333333, 0.666667, 0.166667, 0.5, 0.833333, 0.333333, 0.0, 0.166667, 0.5, 0.333333],
    'PreviousAdInteractions': [2, 1, 3, 0, 2, 1, 0, 1, 2, 1],
    'ShoppingHabits_Frequent': [0, 1, 0, 0, 0, 0, 1, 1, 0, 0],
    'ShoppingHabits_Occasional': [0, 0, 0, 0, 1, 0, 0, 0, 0, 0],
    'ShoppingHabits_RegularBuyer': [1, 0, 0, 1, 0, 1, 0, 0, 0, 0],
    'DeviceType_Desktop': [1, 0, 0, 1, 0, 1, 0, 1, 0, 1],
    'DeviceType_Mobile': [0, 1, 1, 0, 1, 0, 1, 0, 0, 0],
    'Browser_Chrome': [1, 0, 0, 1, 0, 1, 0, 1, 0, 1],
    'Browser_Firefox': [0, 1, 0, 0, 0, 0, 1, 1, 0, 0],
    'Browser_Safari': [0, 0, 1, 0, 1, 0, 0, 0, 1, 0],
    'BannerAd': ['Sports', 'Cinematic', 'Artistic', 'Sports', 'Cinematic', 'Artistic', 'Sports', 'Cinematic',
'Artistic', 'Sports']
})

# Separate the features (X) and target variable (y)
X = data.drop('BannerAd', axis=1)
y = data['BannerAd']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Random Forest classifier
model = RandomForestClassifier()

# Measure the time taken for training
start_time = time.time()
model.fit(X_train, y_train)
training_time = time.time() - start_time

# Make predictions on the test set
y_pred = model.predict(X_test)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)

# Calculate the robustness of the model by performing cross-validation
robustness_scores = cross_val_score(model, X, y, cv=3)  # Adjust cv value as needed
robustness = robustness_scores.mean()

# Calculate the speed of the model by measuring the prediction time
start_time = time.time()
model.predict(X_test.iloc[0:1, :])  # Predict on a single sample for demonstration
prediction_time = time.time() - start_time

# Interpretability of the model is high for Random Forest as it provides feature importances
feature_importances = model.feature_importances_

# Scalability depends on the size of the dataset and the resources available

# Print the evaluation metrics
print("Accuracy:", accuracy)
print("Robustness:", robustness)
print("Training Time:", training_time)
print("Prediction Time:", prediction_time)
print("Feature Importances:", feature_importances)
```

```
Accuracy: 0.5
Robustness: 0.08333333333333333
Training Time: 0.20940518379211426
Prediction Time: 0.022120952606201172
Feature Importances: [0.12935975 0.1313526  0.05869753 0.09065028 0.13563435 0.07663636
 0.08663363 0.03985186 0.02959598 0.04263476 0.02982665 0.02798781
 0.04400823 0.02081597 0.05631423]
```

References

Acar, G., Eubank, C., Englehardt, S., Juarez, M., Narayanan, A., & Diaz, C. (2014, November). The web never forgets: Persistent tracking mechanisms in the wild. In *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security* (pp. 674-689).

Breiman, L. (2001). Random forests. *Machine learning*, *45*, 5-32.

Englehardt, S., & Narayanan, A. (2016, October). Online tracking: A 1-million-site measurement and analysis. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security* (pp. 1388-1401).

Gardner, M. W., & Dorling, S. R. (1998). Artificial neural networks (the multilayer perceptron)—a review of applications in the atmospheric sciences. *Atmospheric environment*, *32*(14-15), 2627-2636.

Hauser, J. R., Liberali, G., & Urban, G. L. (2014). Website morphing 2.0: Switching costs, partial exposure, random exit, and when to morph. *Management science*, *60*(6), 1594-1616.

Kohavi, R. (1995, August). A study of cross-validation and bootstrap for accuracy estimation and model selection. In *Ijcai* (Vol. 14, No. 2, pp. 1137-1145).

Quinlan, J. R. (1986). Induction of decision trees. *Machine learning*, *1*, 81-106.

Rudin, C. (2019). Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature machine intelligence*, *1*(5), 206-215.

Sokolova, M., & Lapalme, G. (2009). A systematic analysis of performance measures for classification tasks. *Information processing & management*, *45*(4), 427-437.

Wolpert, D. H. (1996). The lack of a priori distinctions between learning algorithms. *Neural computation*, *8*(7), 1341-1390.

Xu, C., Liu, D., & Mei, X. (2021). Exploring an efficient POI recommendation model based on user characteristics and spatial-temporal factors. *Mathematics*, *9*(21), 2673.

Keim, D., Andrienko, G., Fekete, J. D., Görg, C., Kohlhammer, J., & Melançon, G. (2008). *Visual analytics: Definition, process, and challenges* (pp. 154-175). Springer Berlin Heidelberg.

Rousseeuw, P. J., & Leroy, A. M. (2005). *Robust regression and outlier detection*. John wiley & sons.

Sterne, J. A., White, I. R., Carlin, J. B., Spratt, M., Royston, P., Kenward, M. G., ... & Carpenter, J. R. (2009). Multiple imputation for missing data in epidemiological and clinical research: potential and pitfalls. *Bmj*, *338*.

Meng, F., Cheng, W., & Wang, J. (2021). Semi-supervised software defect prediction model based on tri-training. *KSII Transactions on Internet & Information Systems*, *15*(11).